

Beli: Discovery Without a Friend Graph


Two interaction-level findings from the restaurant detail page and the map view, with hypotheses on the product tradeoffs behind each.

Context

Beli's core differentiator against Yelp and Google Reviews is that it filters recommendations through your social graph rather than the anonymous crowd. That premise holds well when your friends have already been somewhere. This teardown looks at what happens when they haven't, using Farzi NYC, a restaurant with zero friend activity on my account, and separately examines a map-density issue in the Recs tab.

Finding 1: The friend-coverage gap is a growth lever, not a missing feature


The assumption going in was that Beli simply has nothing to show you when your friend graph hasn't reached a place yet. The screens say otherwise.



Farzi NYC detail page – Rec Score (8.1, based on 1,015 ratings) is shown immediately. Friend Score sits empty beside it.

Beli already has a robust aggregate signal here of 8.1 across 1,015 ratings, and surfaces it as the “Rec Score” without hesitation. So the product is not actually graph-dependent at the top level; it has a perfectly good

fallback to crowd data and uses it. The gap only opens up one layer down, in the “Friend Score” module, which returns a flat dash and a “0” badge with no further information.



Ratings Breakdown is blurred and locked. Ghost bars of a real distribution chart are visible underneath the lock icon.


Scrolling down surfaces the actual mechanic. The “Ratings Breakdown” section, presumably a histogram of how the full rating population scored the place, is blurred out behind a lock icon, with the prompt “Invite a friend to see what all members think of this spot.” The ghost bars underneath the blur are a strong signal that this data already exists and is being withheld for a non-data reason.

This reframes the finding. It isn't “Beli has no fallback for sparse friend graphs.” It's “Beli has a fallback (Rec Score), and then gates a second, more detailed layer of social proof behind an invite action, at the exact moment a user is most curious.” That's a deliberate growth mechanic: the curiosity gap created by a locked breakdown is being used to convert browsing into invites, not absorbed as lost engagement.

HYPOTHESIS

Beli's retention model likely depends more on graph density than on any single session's discovery quality. A user who is mildly inconvenienced by a locked breakdown but invites a friend to unlock it is a better long-term outcome for Beli than a user who gets full crowd data for free and never expands their graph. The product is willing to trade a worse first-time experience at sparse nodes for compounding network growth.

WHAT YOUR FRIENDS THINK – THE SAME PATTERN AT A SMALLER SCALE



“What your friends think” section: zero friends have been here, with a direct “Find friends” CTA.

The “What your friends think” module repeats the same shape: empty state, explicit acknowledgment that friends haven't been there, and an immediate CTA to expand the graph rather than to browse strangers' reviews. There is no option anywhere on this page to see what non-friend members wrote, even in aggregate text form. The numeric Rec Score is the only crowd signal allowed through; written opinions are graph-gated entirely.

THE UX COST: “NO DATA” AND “NO DATA FOR YOU” LOOK IDENTICAL

The one real usability problem here, independent of the growth strategy, is legibility. The Friend Score module renders as a flat dash with a “0” badge, visually the same pattern used elsewhere in the app for an actual score of zero. A user skimming quickly could easily read that as “no one has an opinion on this place,” when the true meaning is “you specifically have zero friends who've been here.” Those are very different situations with very different next actions (one is a dead end, the other has an obvious fix: follow more people). Collapsing them into the same flat-dash treatment undersells the fixable case.


PROPOSED FIX

- Differentiate the empty state visually from a real low score; use an icon or distinct treatment instead of a dash-in-a-circle that mimics the scored state.
- Keep the invite/follow CTA (it's doing real work for growth) but pair it with a lightweight non-graph fallback, such as a one-line aggregate sentiment summary (e.g. “Members love the chaat, mixed on noise level”) so users get partial value before deciding whether to invite anyone.

- Surface the same distribution chart that's already being rendered (just blurred) as a low-resolution preview, for instance by showing the shape of the curve without exact numbers, so the locked state communicates “there's real signal here” rather than reading as a dead module.

Finding 2: The map prioritizes score over status, and breaks down at density

The second issue is a more conventional interaction bug, not a strategic tradeoff.



Recs tab map view in a dense area – overlapping score bubbles with no visual distinction between Been, Want to Try, and Rec pins.

Every pin on this map uses the same visual treatment: a green circle with a fork-and-knife icon, and a number. The only thing that varies is the score. At any meaningful pin density, which is exactly when a map view is most necessary, i.e. dense urban cores with many options nearby, these pins stack and overlap, and several score labels are genuinely unreadable in this screenshot (8.6, 8.7, 9.2, 9.1, 9.0 are all crowded into one cluster).

THE INFORMATION HIERARCHY PROBLEM

Score is the wrong primary visual encoding for a map. In a list view, ranking by score makes sense because you're scanning top to bottom for the best option. On a map, the question is geographic and status-based: “which of these have I been to, which are on my want-to-try list, and which are new recs near me right now?” Status, not score, is the variable that should be encoded most visibly (through color or shape) because it's the one that changes what action a user takes next (revisit vs. explore vs. add to list).

HYPOTHESIS

Beli likely built the map as a secondary view on top of an already-working list/score system, inheriting the score-first visual language by default rather than designing pin encoding specifically for spatial decision-making. This is a common pattern when a list view ships first and a map is added later as a different lens on the same data, without re-examining what the map's primary job is.

PROPOSED FIX

- Encode status (Been / Want to Try / Rec) as pin color or shape, and move score to a secondary label that only appears on tap or at closer zoom.
- Cluster overlapping pins at wide zoom levels into a single marker with a count (e.g. “6 spots”), expanding into individual pins as the user zooms in – standard mapping pattern that Beli doesn't appear to be using here.
- Keep filters (City, Reserve, Order, Open now) accessible from the map but consider a status filter as the default first option, since it's the dimension most map users are trying to parse first.

Summary

Finding	Type	Hypothesis
Locked ratings breakdown	Deliberate growth/monetization tradeoff	Curiosity gap at sparse graph nodes is used to drive invites, prioritizing long-term network density over single-session completeness
Map pin hierarchy	Information-hierarchy bug	List-view visual language (score-first) was inherited by the map without redesigning for spatial, status-based decision-making